

# Project Course 4\_Pokemon\_Final

June 14, 2021

## 1 Pokemon type Clustering

### 1.1 1. Main objectives

The main objective of this work is to clusterise the different types of Pokemons in order to finally observe if one of these kinds present better features than the others. We already have some Pokemon classification as we will see in the data description, but we will drop these corresponding features and try to infer it without knowing.

Data has been obtained from <https://www.kaggle.com/abcsds/pokemon>.

### 1.2 2. Description of the data

The data is imported and the first rows are also shown.

```
[1]: import pandas as pd, seaborn as sns, matplotlib.pyplot as plt, numpy as np

filepath = 'Pokemon.csv'

data = pd.read_csv(filepath, sep = ',')

data.head()
```

```
[1]:   number          name  type1  type2  total  hp  attack  defense \
0       1      Bulbasaur  Grass  Poison    318  45      49       49
1       2      Ivysaur   Grass  Poison    405  60      62       63
2       3      Venusaur  Grass  Poison    525  80      82       83
3       3  Mega Venusaur  Grass  Poison    625  80     100      123
4       3  Gigantamax Venusaur  Grass  Poison    525  80      82       83

   sp_attack  sp_defense  speed  generation  legendary
0        65         65     45           1     False
1        80         80     60           1     False
2       100        100     80           1     False
3       122        120     80           1     False
4       100        100     80           1     False
```

Each column corresponds to the following:

**Number:** The ID for each pokemon

**Name:** The name of each pokemon

**Type 1:** Each pokemon has a type, this determines weakness/resistance to attacks

**Type 2:** Some pokemon are dual type and have 2

**Total:** Sum of all stats that come after this, a general guide to how strong a pokemon is

**HP:** Hit points, or health, defines how much damage a pokemon can withstand before fainting

**Attack:** The base modifier for normal attacks (eg. Scratch, Punch)

**Defense:** The base damage resistance against normal attacks

**SP Atk:** Special attack, the base modifier for special attacks (e.g. fire blast, bubble beam)

**SP Def:** Special defense, the base damage resistance against special attacks

**Speed:** Determines which pokemon attacks first each round

**Generation:** The generation of games where the pokemon was first introduced

**Legendary:** Some pokemon are much rarer than others, and are dubbed “legendary”

We have 13 features and 1072 rows. Let's also take some extra information. For the *type2* feature he have a large number of null values.

[2]: `data.shape`

[2]: `(1072, 13)`

[3]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1072 entries, 0 to 1071
Data columns (total 13 columns):
 #   Column      Non-Null Count  Dtype  
 ---  --          -----          ----  
 0   number      1072 non-null   int64  
 1   name        1072 non-null   object  
 2   type1       1072 non-null   object  
 3   type2       574 non-null    object  
 4   total        1072 non-null   int64  
 5   hp          1072 non-null   int64  
 6   attack       1072 non-null   int64  
 7   defense      1072 non-null   int64  
 8   sp_attack    1072 non-null   int64  
 9   sp_defense   1072 non-null   int64  
 10  speed        1072 non-null   int64  
 11  generation   1072 non-null   int64  
 12  legendary     1072 non-null   bool
```

```
dtypes: bool(1), int64(9), object(3)
memory usage: 101.7+ KB
```

### 1.3 3. Feature engineering

In order to do clustering, we will remove the **type1**, **type2**, **name**, **numberID**, **legendary** and **generation**. Except *name* and *numberID*, they are categorical data, and we are not allowed to work with them. We will only work with continuous data.

```
[7]: data = data.drop(['name', 'number', 'type1', 'type2', 'legendary', ↴'generation'], axis = 1)
```

```
[8]: data
```

```
[8]:      total    hp   attack  defense  sp_attack  sp_defense  speed
 0       318    45      49      49        65         65       45
 1       405    60      62      63        80         80       60
 2       525    80      82      83       100        100       80
 3       625    80     100     123       122        120       80
 4       525    80      82      83       100        100       80
 ...
 ...
 1067    580   100     145     130        65        110       30
 1068    580   100      65      60       145         80      130
 1069    500   100      80      80        80         80       80
 1070    680   100     165     150        85        130       50
 1071    680   100      85      80       165        100      150
[1072 rows x 7 columns]
```

Let's first take a look to the correlations:

```
[9]: corr_mat = data.corr()

for x in range(len(data.columns)):
    corr_mat.iloc[x,x] = 0

corr_mat.abs().idxmax()
```

```
[9]: total          attack
      hp            total
      attack        total
      defense       total
      sp_attack     total
      sp_defense    total
      speed         total
      dtype: object
```

We observe that the general indicator of how strong a pokemon is (*total*) is highly correlated with the power *attack*.

Now, the **skew of each feature is evaluated**, and we will correct it by applying the logarithm for those who have skew > 0.75.

```
[10]: skew_columns = data.skew().sort_values(ascending = False)
```

```
skew_columns = skew_columns.loc[skew_columns > 0.75]  
skew_columns
```

```
[10]: hp           1.760494  
defense       1.143146  
sp_defense    0.926515  
dtype: float64
```

```
[11]: for col in skew_columns.index:  
      data[col] = np.log1p(data[col])
```


Now all the features are reescaled using the **MinMaxScaler**.

```
[12]: from sklearn.preprocessing import MinMaxScaler  
  
float_cols = [x for x in data.columns]  
  
mm = MinMaxScaler()  
  
data[float_cols] = mm.fit_transform(data[float_cols])
```

Once we have done with the sclaning, we can check the correlation with a pairplot. We will also differentiate by the **legendary** feature, to see if there is a kind of clustering. The **type1** and **type2** features have also been checked before and it is observed a random distribution in the correlation plot.

```
[13]: sns.set_context('notebook')  
  
data['legendary'] = df['legendary']  
sns.pairplot(data[float_cols + ['legendary']], hue = 'legendary')
```

```
[13]: <seaborn.axisgrid.PairGrid at 0x7f8b192f16d0>
```



A kind of clustering is observed for the **legendary** feature.

```
[14]: data = data.drop('legendary', axis = 1)
```

```
[15]: data
```

	total	hp	attack	defense	sp_attack	sp_defense	speed
0	0.150526	0.646223	0.237838	0.567873	0.298913	0.461574	0.205128
1	0.242105	0.704391	0.308108	0.633990	0.380435	0.544121	0.282051
2	0.368421	0.762836	0.416216	0.706822	0.489130	0.633068	0.384615
3	0.473684	0.762836	0.513514	0.811133	0.608696	0.705892	0.384615
4	0.368421	0.762836	0.416216	0.706822	0.489130	0.633068	0.384615
...	...	...	...	...	...	...	...
1067	0.426316	0.808316	0.756757	0.825841	0.298913	0.671122	0.128205
1068	0.426316	0.808316	0.324324	0.621131	0.733696	0.544121	0.641026

```
1069  0.342105  0.808316  0.405405  0.697082  0.380435  0.544121  0.384615  
1070  0.531579  0.808316  0.864865  0.863895  0.407609  0.737898  0.230769  
1071  0.531579  0.808316  0.432432  0.697082  0.842391  0.633068  0.743590
```

```
[1072 rows x 7 columns]
```

## 1.4 4. Three variations of unsupervised model (Clustering)

In this section we apply 3 different models of clustering. After that we will evaluate which of them performs better.

### 1.4.1 4.1 K-means


The **K-means** model is applied for the different number of clusters. For each one of them the inertia is computed, and at this point the elbow rule is used in order to take the best number of **k**.

```
[16]: from sklearn.cluster import KMeans  
km_list = list()  
  
for clust in range(1,21):  
    km = KMeans(n_clusters=clust, random_state= 42)  
    km = km.fit(data)  
    km_list.append(pd.Series({'clusters': clust, 'inertia':km.inertia_, 'model':  
    ↴ km}))
```

```
[17]: plot_data = pd.concat(km_list, axis = 1).T  
plot_data = plot_data[['clusters', 'inertia']].set_index('clusters')  
plot_data.head(4)
```

```
[17]:      inertia  
clusters  
1        156.274  
2        97.0283  
3        80.7911  
4        70.9049
```

```
[18]: ax = plot_data.plot(marker = 'o', color = 'blue')  
ax.set_xticks(range(0,21,2))  
ax.set_xlim(0,21)  
ax.set_xlabel('Cluster', ylabel='Inertia');
```




We will take  $k = 4$ .

```
[19]: km = KMeans(n_clusters=4, random_state= 42)
km = km.fit(data[float_cols])
data['kmeans'] = km.predict(data[float_cols])
```

```
[20]: sns.scatterplot(data=data, x="total", y="attack", hue="kmeans")
```

```
[20]: <AxesSubplot:xlabel='total', ylabel='attack'>
```



The algorithm differentiate 4 considerably good clusters. Visually we might only have taken 2 clusters, but we really do not know how many clusters there are in this data.

#### 1.4.2 4.2 Agglomerative Clustering

Now the Agglomerative Clustering is applied for the same number of clusters  $k = 4$ .


```
[21]: from sklearn.cluster import AgglomerativeClustering

agg = AgglomerativeClustering(n_clusters=4, linkage='ward',
                                compute_full_tree=True)
agg = agg.fit(data[float_cols])

data['agg'] = agg.fit_predict(data[float_cols])
```

```
[22]: sns.scatterplot(data=data, x="total", y="attack", hue="agg")
```

```
[22]: <AxesSubplot:xlabel='total', ylabel='attack'>
```



The cluster differentiation is not that clear.

#### 1.4.3 4.3 DBSCAN

DBSCAN is applied for k=4

```
[23]: from sklearn.cluster import DBSCAN
from sklearn import metrics
from sklearn.datasets import make_blobs
from sklearn.preprocessing import StandardScaler

db = DBSCAN(eps=0.2, min_samples=15).fit(data[float_cols])
clustering_labels = db.fit_predict(data[float_cols])


data['dbscan'] = clustering_labels
```

```
[24]: data['dbscan'].value_counts()
```

```
[24]: 0    986
-1    86
Name: dbscan, dtype: int64
```

```
[25]: sns.scatterplot(data=data, x="total", y="attack", hue="dbscan")
```

```
[25]: <AxesSubplot:xlabel='total', ylabel='attack'>
```



Choosing the parameters is really complicated. Nonetheless, visually we could see that at least we have 2 clusters, but DBSCAN is not capable of determining them well, this is because **we have different cluster densities**.

## 1.5 5. Which model is recommended?

In this case I would recommend **k-means**. It is capable of differentiating clusters quite well, and with different cluster densities.

The DBSCAN is complicated to apply since density of points is not uniform, and choosing parameters is not trivial either. Very quickly we obtain a large number of clusters, or directly simply one just like I have shown in the last figure.

## 1.6 6. Future suggestions

In the future maybe different distances for the Agglomerative Clustering method could be checked, and also a grid cross validation for parameters of DSCAN as well. Nonetheless, in this example not very good results would have been obtained.

```
[ ]:
```